

Other Relational Languages

Tuple Relational Calculus

- A nonprocedural query language, where each query is of the form

$$\{t \mid P(t)\}$$

- It is the set of all tuples t such that predicate P is true for t
- t is a *tuple variable*, $t[A]$ denotes the value of tuple t on attribute A
- $t \in r$ denotes that tuple t is in relation r
- P is a *formula* similar to that of the predicate calculus

Predicate Calculus Formula

1. Set of attributes and constants
2. Set of comparison operators: (e.g., $<$, \leq , $=$, \neq , $>$, \geq)
3. Set of connectives: and (\wedge), or (\vee), not (\neg)
4. Implication (\Rightarrow): $x \Rightarrow y$, if x is true, then y is true

$$x \Rightarrow y \equiv \neg x \vee y$$

5. Set of quantifiers:

- ▶ Existential quantifier:

- ▶ $\exists t \in r(Q(t)) \equiv$ "there exists" a tuple t in relation r such that predicate $Q(t)$ is true

- ▶ Universal quantifier:

- ▶ $\forall t \in r(Q(t)) \equiv Q$ is true "for all" tuples t in relation r

Example Queries

- University Example:

- `instructor(ID, name, dept_name, salary)`
- `section(course_id, semester, year, sec_id)`

- Find the *ID*, *name*, *dept_name*, *salary* for instructors whose salary is greater than \$80,000

$$\{t \mid t \in \text{instructor} \wedge t[\text{salary}] > 80000\}$$

- As in the previous query, but output only the *ID* attribute value

$$\{t \mid \exists s \in \text{instructor} (t[ID] = s[ID] \wedge s[\text{salary}] > 80000)\}$$

Notice that a relation on schema (*ID*) is *implicitly* defined by the query

Example Queries

- Find the names of all instructors whose department is in the Watson building

$$\{t \mid \exists s \in \text{instructor} (t[\text{name}] = s[\text{name}] \wedge \exists u \in \text{department} (u[\text{dept_name}] = s[\text{dept_name}] \wedge u[\text{building}] = \text{"Watson"}))\}$$

- Find the set of all courses taught in the Fall 2009 semester, or in the Spring 2010 semester, or both

$$\{t \mid \exists s \in \text{section} (t[\text{course_id}] = s[\text{course_id}] \wedge s[\text{semester}] = \text{"Fall"} \wedge s[\text{year}] = 2009 \vee \exists u \in \text{section} (t[\text{course_id}] = u[\text{course_id}] \wedge u[\text{semester}] = \text{"Spring"} \wedge u[\text{year}] = 2010))\}$$

Example Queries

- Find the set of all courses taught in the Fall 2009 semester, and in the Spring 2010 semester

$$\{t \mid \exists s \in \text{section} (t[\text{course_id}] = s[\text{course_id}] \wedge s[\text{semester}] = \text{"Fall"} \wedge s[\text{year}] = 2009 \wedge \exists u \in \text{section} (t[\text{course_id}] = u[\text{course_id}] \wedge u[\text{semester}] = \text{"Spring"} \wedge u[\text{year}] = 2010))\}$$

- Find the set of all courses taught in the Fall 2009 semester, but not in the Spring 2010 semester

$$\{t \mid \exists s \in \text{section} (t[\text{course_id}] = s[\text{course_id}] \wedge s[\text{semester}] = \text{"Fall"} \wedge s[\text{year}] = 2009 \wedge \neg \exists u \in \text{section} (t[\text{course_id}] = u[\text{course_id}] \wedge u[\text{semester}] = \text{"Spring"} \wedge u[\text{year}] = 2010))\}$$

Safety of Expressions

- It is possible to write tuple calculus expressions that generate infinite relations.
- For example, $\{ t \mid \neg t \in r \}$ results in an infinite relation if the domain of any attribute of relation r is infinite
- To guard against the problem, we restrict the set of allowable expressions to safe expressions.
- An expression $\{ t \mid P(t) \}$ in the tuple relational calculus is *safe* if every component of t appears in **dom**(P) : the relations, tuples, or constants that appear in P
 - NOTE: this is more than just a syntax condition.
 - ▶ E.g. $\{ t \mid t[A] = 5 \vee \mathbf{true} \}$ is not safe --- it defines an infinite set with attribute values that do not appear in any relation or tuples or constants in P .

Universal Quantification

- Find all students who have taken all courses offered in the Biology department, and the courses they took:

- $\{t \mid \forall u \in \text{course } (u[\text{dept_name}] = \text{"Biology"} \Rightarrow$
 $\exists s \in \text{takes } (t[ID] = s[ID] \wedge$
 $s[\text{course_id}] = u[\text{course_id}] \wedge$
 $t[\text{course_id}] = s[\text{course_id}])\}$
- Note the schema on r is $(ID, \text{course_id})$ here.
- What's the problem with this query?
- The above query would be unsafe if the Biology department has not offered any courses.

Universal Quantification

- Find all students who have taken all courses offered in the Biology department, **and the courses they took:**

- $\{t \mid \exists r \in \text{student} (t[ID] = r[ID]) \wedge$
 $(\forall u \in \text{course} (u[\text{dept_name}] = \text{"Biology"} \Rightarrow$
 $\exists s \in \text{takes} (t[ID] = s[ID] \wedge$
 $s[\text{course_id}] = u[\text{course_id}] \wedge$
 $t[\text{course_id}] = s[\text{course_id}]))\}$
- Add the **existential quantification** on student.

Domain Relational Calculus

- A nonprocedural query language equivalent in power to the tuple relational calculus
- Each query is an expression of the form:

$$\{ \langle x_1, x_2, \dots, x_n \rangle \mid P(x_1, x_2, \dots, x_n) \}$$

- x_1, x_2, \dots, x_n represent domain variables
- Gives the schema of the output relation explicitly
- P represents a formula similar to that of the predicate calculus

Example Queries

- Find the *ID*, *name*, *dept_name*, *salary* for instructors whose salary is greater than \$80,000
 - $\{ \langle i, n, d, s \rangle \mid \langle i, n, d, s \rangle \in instructor \wedge s > 80000 \}$
- As in the previous query, but output only the *ID* attribute value
 - $\{ \langle i \rangle \mid \exists n, d, s (\langle i, n, d, s \rangle \in instructor \wedge s > 80000) \}$
- Find the names of all instructors whose department is in the Watson building
 - $\{ \langle n \rangle \mid \exists i, d, s (\langle i, n, d, s \rangle \in instructor \wedge \exists b, a (\langle d, b, a \rangle \in department \wedge b = \text{“Watson”})) \}$

Example Queries

- Find the set of all courses taught in the Fall 2009 semester, or in the Spring 2010 semester, or both

$$\{ \langle c \rangle \mid \exists a, s, y, b, t (\langle c, a, s, y, b, t \rangle \in \text{section} \wedge s = \text{"Fall"} \wedge y = 2009) \vee \exists a, s, y, b, t (\langle c, a, s, y, b, t \rangle \in \text{section} \wedge s = \text{"Spring"} \wedge y = 2010) \}$$

This case can also be written as

$$\{ \langle c \rangle \mid \exists a, s, y, b, t (\langle c, a, s, y, b, t \rangle \in \text{section} \wedge ((s = \text{"Fall"} \wedge y = 2009) \vee (s = \text{"Spring"} \wedge y = 2010))) \}$$

- Find the set of all courses taught in the Fall 2009 semester, and in the Spring 2010 semester

$$\{ \langle c \rangle \mid \exists a, s, y, b, t (\langle c, a, s, y, b, t \rangle \in \text{section} \wedge s = \text{"Fall"} \wedge y = 2009) \wedge \exists a, s, y, b, t (\langle c, a, s, y, b, t \rangle \in \text{section} \wedge s = \text{"Spring"} \wedge y = 2010) \}$$

Safety of Expressions

The expression:

$$\{ \langle x_1, x_2, \dots, x_n \rangle \mid P(x_1, x_2, \dots, x_n) \}$$

is safe if all of the following hold:

1. All values that appear in tuples of the expression are values from *dom*(*P*) (that is, the values appear either in *P* or in a tuple of a relation mentioned in *P*).
2. For every “there exists” subformula of the form $\exists x (P_1(x))$, the subformula is true if and only if there is a value of *x* in *dom*(*P*₁) such that *P*₁(*x*) is true.
3. For every “for all” subformula of the form $\forall x (P_1(x))$, the subformula is true if and only if *P*₁(*x*) is true for all values *x* from *dom*(*P*₁).

Universal Quantification

- Find all students who have taken all courses offered in the Biology department
 - $\{ \langle i \rangle \mid \exists n, d, tc (\langle i, n, d, tc \rangle \in student \wedge$
 $(\forall ci, ti, dn, cr (\langle ci, ti, dn, cr \rangle \in course \wedge dn = \text{“Biology”}$
 $\Rightarrow \exists si, se, y, g (\langle i, ci, si, se, y, g \rangle \in takes))) \}$
 - Note that without the existential quantification on student, the above query would be unsafe if the Biology department has not offered any courses.

Datalog

- Non-procedural query language based on Prolog.
- A Datalog program consists of a set of **rules**, each defines a *view*.

- `v1(A, B) :- account(A, "Perryridge", B), B > 700.`

- Commas “,” read as “AND”.

Body

Head ■ To retrieve the balance of account A-314:

- `?- v1("A-314", B).`
- Answer: ("A-314", 780).

- To get the account number and balance of all accounts in v1 with balance more than 1000:

- `?- v1(A, B), B > 1000.`
- Answer: ("A-205", 1200).

Syntax of Datalog Rules

- Uppercase letters or words starting with uppercase letters as variables
- Lowercase letters and words starting with lowercase letter as relation names and attribute names.
- Positive literal: $p(t_1, t_2, \dots, t_n)$
 - t_1, t_2, \dots, t_n are either constants or variables.
 - p is the predicate symbol.
- Negative literal: **not** $p(t_1, t_2, \dots, t_n)$
- $B > 700$ can be understood as a literal, too: $> (B, 700)$
- $p(v_1, v_2, \dots, v_n)$ is a fact, where v_1, \dots, v_n are constants.
 - Tuple (v_1, v_2, \dots, v_n) is in relation p .
- A rule is expressed as:

$$p(t_1, t_2, \dots, t_n) \text{ :- } L_1, L_2, \dots, L_m$$

$(L_1, L_2, \dots, L_m \text{ are literals.})$

Semantic of Rules

- The set of facts that can be inferred from a given set of facts I , given a rule R :
 - $\text{infer}(R, I) = \{p(t1, \dots, tn) \mid \text{there is an instantiation } R' \text{ of } R, \text{ where } p(t1, \dots, tn) \text{ is the head of } R', \text{ and the body of } R' \text{ is satisfied by } I.\}$
- Given a set of rules $\mathcal{R} = \{R1, R2, \dots, Rn\}$,
 - $\text{infer}(\mathcal{R}, I) = \text{infer}(R1, I) \cup \text{infer}(R2, I) \cup \dots \cup \text{infer}(Rn, I)$
 - This set of rules is basically the Datalog program
- Non-recursive Datalog without arithmetics have equivalent expressive power to basic relational algebra
 - Try this: Find the set of all courses taught in the Fall 2009 semester, but not in the Spring 2010 semester?

Recursion in Datalog

- Deals with recursive data structure, e.g. lists and trees.
- Report_Schema = (employee_name, manager_name)
- Supposed we want to find all employees managed by a person X.

managed_by (Y, X) :- report (Y, X).

managed_by (Y, X) :- report(Y, Z), managed_by (Z, X).

- Recursive Datalog contains no negative literals.
- Recursive rules are evaluated by *iteratively* computing the **fix point** or the condition under which no new facts can be inferred → termination.

E-R Model (I)

Modeling

- A *database* can be modeled as:
 - a collection of entities,
 - relationship among entities.
- An **entity** is an object that exists and is distinguishable from other objects.
 - Example: specific person, company, event, plant
- Entities have **attributes**
 - Example: people have *names* and *addresses*
- An **entity set** is a set of entities of the same type that share the same properties.
 - Example: set of all persons, companies, trees, holidays

Entity Sets *instructor* and *student*

instructor_ID instructor_name

76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

instructor

student-ID student_name

98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

student

Relationship Sets

- A **relationship** is an association among several (typically two) entities

Example:

44553 (Peltier)
student entity

advisor
relationship set

22222 (Einstein)
instructor entity

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$


where (e_1, e_2, \dots, e_n) is a relationship

- Example:

$$(44553, 22222) \in \text{advisor}$$

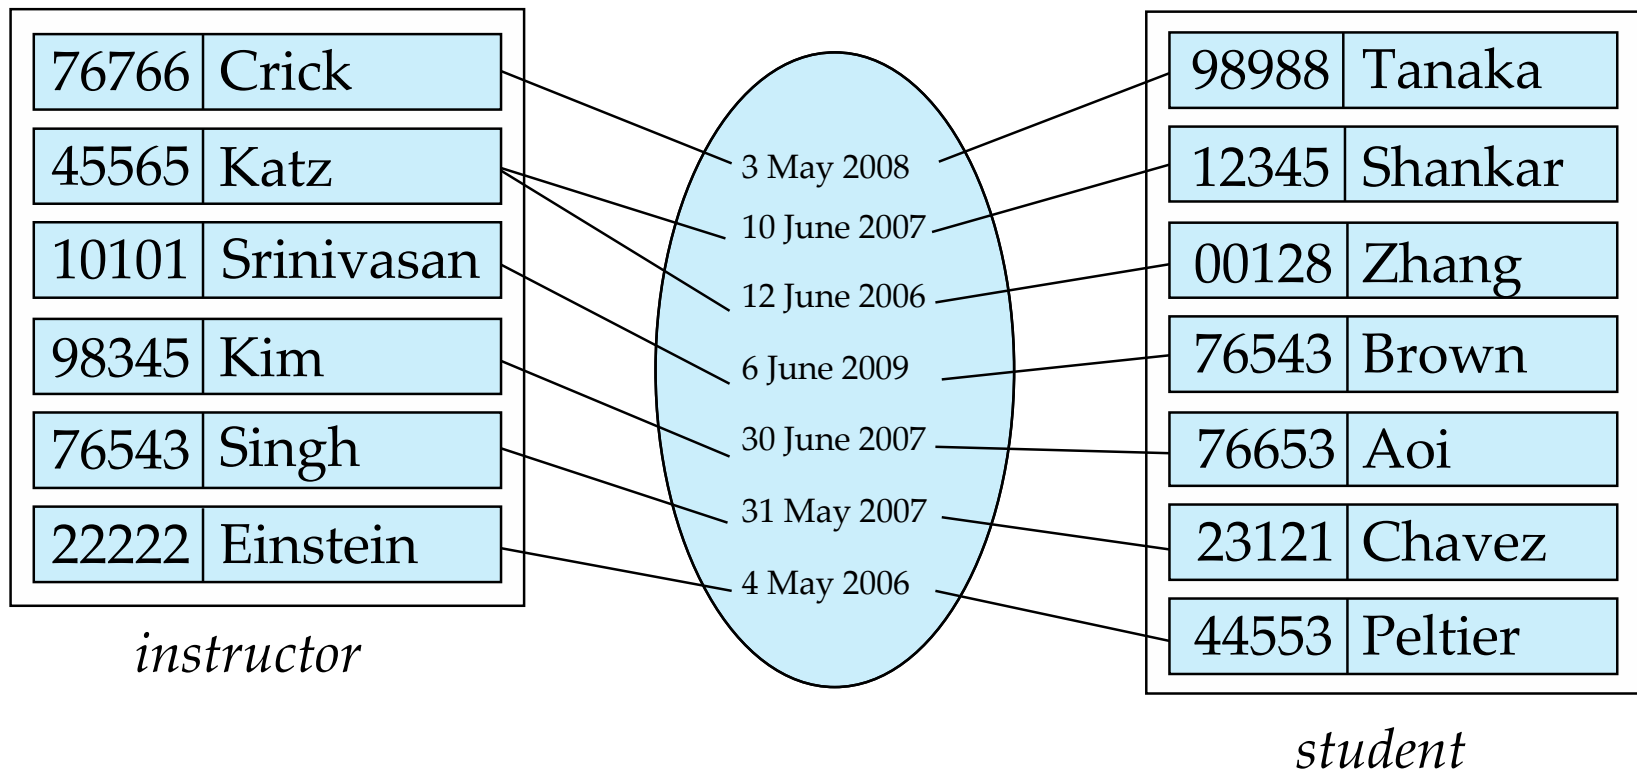
- Basically, a *relationship* is a tuple, while a *relationship set* is a set of tuples

Relationship Set *advisor*



Relationship Sets (Cont.)

- An **attribute** can also be property of a relationship set.
- For instance, the *advisor* relationship set between entity sets *instructor* and *student* may have the attribute *date* which tracks when the student started being associated with the advisor



Degree of a Relationship Set

■ binary relationship

- involve two entity sets (or degree two).
- most relationship sets in a database system are binary.

■ Relationships between more than two entity sets are rare. Most relationships are binary. (More on this later.)

- ▶ Example: *students* work on research *projects* under the guidance of an *instructor*.
- ▶ relationship *proj_guide* is a ternary relationship between *instructor*, *student*, and *project*

Attributes

- An entity is represented by a set of attributes, that is descriptive properties possessed by all members of an entity set.

- Example:

instructor = (ID, name, street, city, salary)

course= (course_id, title, credits)


- **Domain** – the set of permitted values for each attribute

- Attribute types:


- **Simple** and **composite** attributes.
- **Single-valued** and **multivalued** attributes
 - ▶ Example: multivalued attribute: *phone_numbers*
- **Derived** attributes
 - ▶ Can be computed from other attributes
 - ▶ Example: age, given date_of_birth

Composite Attributes

composite
attributes




component
attributes




Mapping Cardinality Constraints

- Express the number of entities to which another entity can be associated via a relationship set.
- Most useful in describing binary relationship sets.
- For a binary relationship set the mapping cardinality must be one of the following types:
 - One to one
 - One to many
 - Many to one
 - Many to many

Mapping Cardinalities




One to one



One to many


Note: Some elements in A and B may not be mapped to any elements in the other set (partial mapping)

Mapping Cardinalities



(a)

Many to
one



(b)

Many to many

Note: Some elements in A and B may not be mapped to any elements in the other set (partial mapping)

Keys

- A **super key** of an entity set is a set of one or more attributes whose values uniquely determine each entity.
- A **candidate key** of an entity set is a minimal super key
 - *ID* is candidate key of *instructor*
 - *course_id* is candidate key of *course*
- Although several candidate keys may exist, one of the candidate keys is selected to be the **primary key**.


Keys for Relationship Sets

- The combination of primary keys of the participating entity sets forms a super key of a relationship set.
 - (s_id, i_id) is the super key of *advisor*
 - *NOTE: this means **a pair of entities can have at most one relationship in a particular relationship set.***
 - ▶ Example: if we wish to track multiple meeting dates between a student and her advisor, we cannot assume a relationship for each meeting. We can use a multivalued attribute though.
- Must consider the mapping cardinality of the relationship set when deciding what are the candidate keys
- Need to consider semantics of relationship set in selecting the *primary key* in case of more than one candidate key

Redundant Attributes

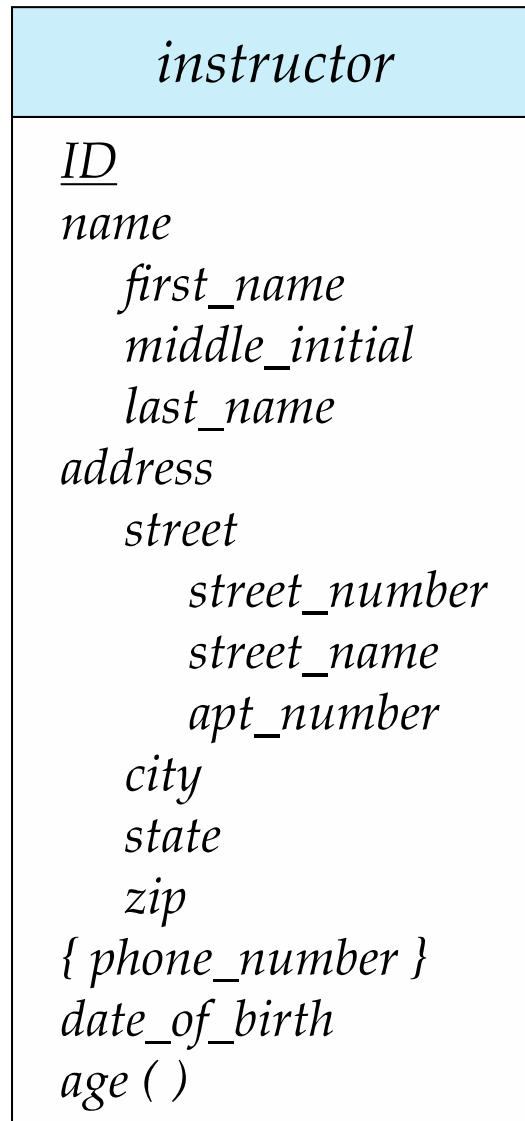
- Suppose we have entity sets
 - *instructor*, with attributes including *dept_name*
 - *department*and a relationship
 - *inst_dept* relating *instructor* and *department*
- Attribute *dept_name* in entity *instructor* is redundant since there is an explicit relationship *inst_dept* which relates instructors to departments
 - The attribute replicates information present in the relationship, and should be removed from *instructor*
 - BUT: when converting back to tables, in some cases the attribute gets reintroduced, as we will see.

E-R Diagrams




- Rectangles represent entity sets.
- Diamonds represent relationship sets.
- Attributes listed inside entity rectangle
- Underline indicates primary key attributes
- Note: we use a slight different (and simplified) notation for entity sets and attributes here!

Entity With Composite, Multivalued, and Derived Attributes

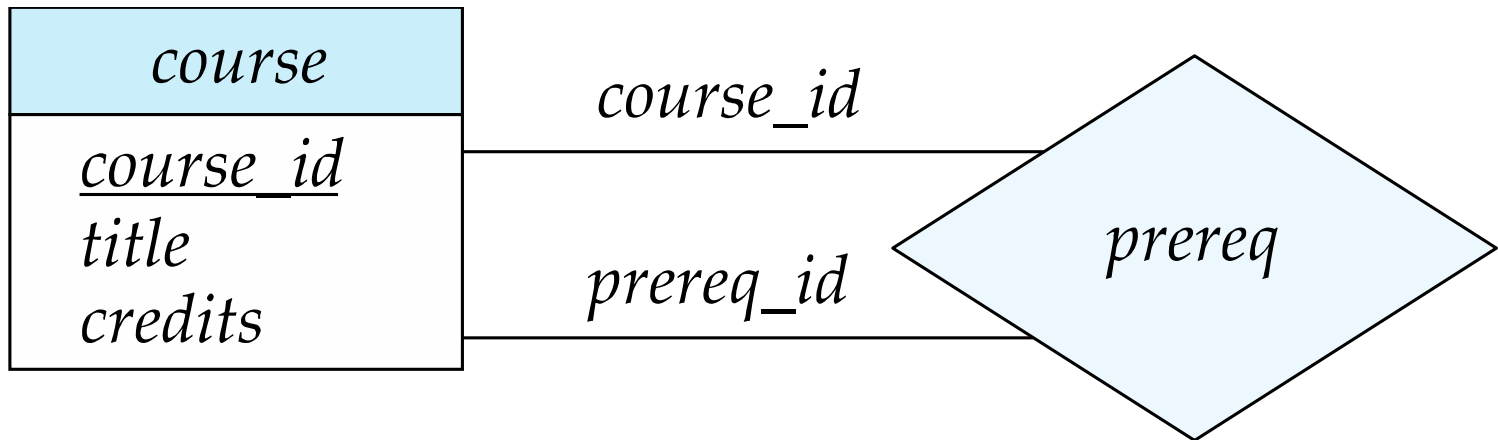


Relationship Sets with Attributes



Roles

- Entity sets of a relationship need not be distinct
 - Each occurrence of an entity set plays a “role” in the relationship
- The labels “*course_id*” and “*prereq_id*” are called **roles**.




Cardinality Constraints

- We express cardinality constraints by drawing either a directed line (\rightarrow), signifying “one,” or an undirected line ($—$), signifying “many,” between the relationship set and the entity set.
- One-to-one relationship:
 - A student is associated with at most one *instructor* via the relationship *advisor*
 - A *student* is associated with at most one *department* via *stud_dept*


One-to-One Relationship

- one-to-one relationship between an *instructor* and a *student*
 - an instructor is associated with at most one student via *advisor*
 - and a student is associated with at most one instructor via *advisor*




One-to-Many Relationship

- one-to-many relationship between an *instructor* and a *student*
 - an instructor is associated with several (including 0) students via *advisor*
 - a student is associated with at most one instructor via *advisor*,




Many-to-One Relationships

- In a many-to-one relationship between an *instructor* and a *student*,
 - an instructor is associated with at most one student via *advisor*,
 - and a student is associated with several (including 0) instructors via *advisor*




Many-to-Many Relationship

- An instructor is associated with several (possibly 0) students via *advisor*
- A student is associated with several (possibly 0) instructors via *advisor*



Participation of an Entity Set in a Relationship Set

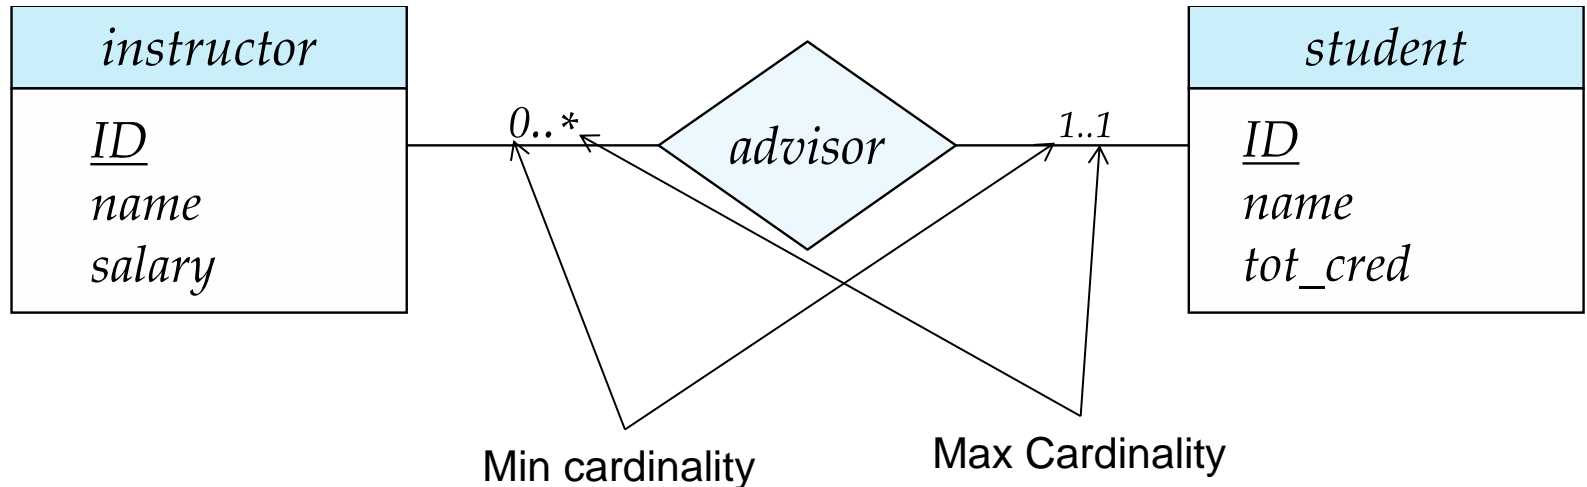
- Total participation (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set
 - E.g., participation of *section* in *sec_course* is total
 - ▶ every *section* must have an associated course
- Partial participation: some entities may not participate in any relationship in the relationship set
 - Example: participation of *instructor* in *advisor* is partial




Note: doubly outlined diamond is a relationship set that identifies a **weak entity set**

Alternative Notation for Cardinality Limits

- Cardinality limits can also express participation constraints



E-R Diagram with a Ternary Relationship



Cardinality Constraints on Ternary Relationship

- We allow at most one arrow out of a ternary (or greater degree) relationship to indicate a cardinality constraint
- E.g., an arrow from *proj_guide* to *instructor* indicates each student has at most one guide for a project
- If there is more than one arrow, there are two ways of defining the meaning.
 - E.g., a ternary relationship R between A , B and C with arrows to B and C could mean
 1. each A entity is associated with a unique entity from B and C or
 2. each pair of entities from (A, B) is associated with a unique C entity, and each pair (A, C) is associated with a unique B
 - Each alternative has been used in different formalisms
 - To avoid confusion we outlaw more than one arrow

End